

## Import Libraries & Load Data

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import confusion_matrix, classification_report,
accuracy_score, roc_auc_score, roc_curve
import warnings
warnings.filterwarnings('ignore')

# Load data
df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")
df.head()

    customerID  gender  SeniorCitizen Partner Dependents  tenure
PhoneService \
0  7590-VHVEG  Female          0      Yes        No       1
No
1  5575-GNVDE   Male          0      No        No      34
Yes
2  3668-QPYBK   Male          0      No        No       2
Yes
3  7795-CFOCW   Male          0      No        No      45
No
4  9237-HQITU  Female          0      No        No       2
Yes

    MultipleLines InternetService OnlineSecurity ...
DeviceProtection \
0  No phone service           DSL        No ...
No
1                  No           DSL        Yes ...
Yes
2                  No           DSL        Yes ...
No
3  No phone service           DSL        Yes ...
Yes
4                  No     Fiber optic        No ...
No

    TechSupport StreamingTV StreamingMovies          Contract
PaperlessBilling \
0            No          No        No Month-to-month
Yes
1            No          No        No      One year
No
```

2	No	No	No	Month-to-month
Yes				
3	Yes	No	No	One year
No				
4	No	No	No	Month-to-month
Yes				

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

## Data Cleaning

```
# Check info
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   customerID      7043 non-null   object 
 1   gender          7043 non-null   object 
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object 
 4   Dependents     7043 non-null   object 
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object 
 7   MultipleLines   7043 non-null   object 
 8   InternetService 7043 non-null   object 
 9   OnlineSecurity  7043 non-null   object 
 10  OnlineBackup    7043 non-null   object 
 11  DeviceProtection 7043 non-null   object 
 12  TechSupport     7043 non-null   object 
 13  StreamingTV     7043 non-null   object 
 14  StreamingMovies  7043 non-null   object 
 15  Contract        7043 non-null   object 
 16  PaperlessBilling 7043 non-null   object 
 17  PaymentMethod    7043 non-null   object 
 18  MonthlyCharges  7043 non-null   float64
 19  TotalCharges    7043 non-null   object 
 20  Churn           7043 non-null   object 
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```

# Convert TotalCharges to numeric
df['TotalCharges']=pd.to_numeric(df['TotalCharges'],errors='coerce')
'''errors : {'ignore', 'raise', 'coerce'}, default 'raise'
- If 'raise', then invalid parsing will raise an exception.
- If 'coerce', then invalid parsing will be set as NaN.
- If 'ignore', then invalid parsing will return the input'''

"errors : {'ignore', 'raise', 'coerce'}, default 'raise'\n      - If
'raise', then invalid parsing will raise an exception.\n      - If
'coerce', then invalid parsing will be set as NaN.\n      - If 'ignore',
then invalid parsing will return the input"

#check null value
df.isnull().sum()

customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents     0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    11
Churn           0
dtype: int64

#fill null value with median
df['TotalCharges']=df['TotalCharges'].fillna(df['TotalCharges'].median())


# Drop customerID
df=df.drop('customerID',axis=1)
#axis : {0 or 'index', 1 or 'columns'}, default 0

# change target variable in binary format
df['Churn']=df['Churn'].map({'Yes':1,'No':0})

```

## Exploratory Data Analysis (EDA)

```
# Check churn balance
ax=sns.countplot(x='Churn',data=df)
ax.bar_label(ax.containers[0])
plt.title('Churn Distribution')
plt.show()
```




```
chrn_count=df['Churn'].value_counts()
chrn_count

Churn
0    5174
1    1869
Name: count, dtype: int64

plt.figure(figsize=(6,6))
plt.pie(chrn_count, labels=['No Churn', 'Churn'], autopct='%.1f%%')
plt.title('Percentage Of Churn Customer')


Text(0.5, 1.0, 'Percentage Of Churn Customer')
```

Percentage Of Churn Customer



in above we understand 73% people is not chrn and 27% people is chrun out

```
plt.figure(figsize=(9,4))
sns.histplot(x='tenure',data=df,bins=72,hue='Churn')
plt.show()
```



in above we understand who have used services is long time its stay who using only 1 and 2 month is churned

```

df.columns

Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
       'PhoneService', 'MultipleLines', 'InternetService',
       'OnlineSecurity',
       'OnlineBackup', 'DeviceProtection', 'TechSupport',
       'StreamingTV',
       'StreamingMovies', 'Contract', 'PaperlessBilling',
       'PaymentMethod',
       'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')

# Columns to plot
cols = [
    'PhoneService', 'MultipleLines', 'InternetService',
    'OnlineSecurity',
    'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
    'StreamingMovies'
]

# Create subplot grid (3 rows x 3 columns)
fig, axes = plt.subplots(3, 3, figsize=(15, 12))
axes = axes.flatten() # Flatten 2D array of axes for easy iteration

# Loop through columns
for i, col in enumerate(cols):
    sns.countplot(x=col, data=df, hue='Churn', ax=axes[i])
    axes[i].set_title(f'{col} vs Churn', fontsize=12)
    axes[i].bar_label(axes[i].containers[0], fontsize=8)

```


```

if len(axes[i].containers) > 1:
    axes[i].bar_label(axes[i].containers[1], fontsize=8)
axes[i].set_xlabel('')
axes[i].set_ylabel('Count')

# Remove extra empty plots (if any)
for j in range(len(cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



in above Customers without internet-based add-ons (like OnlineSecurity, TechSupport, or DeviceProtection) tend to churn more frequently. Fiber optic users also exhibit higher churn rates compared to DSL users. Overall, having additional services (security, backup, or protection) is linked with lower churn

```


# analysis using paymnnet method
plt.figure(figsize=(8,4))

```

```

ax=sns.countplot(x='PaymentMethod',hue='Churn',data=df)
for containers in ax.containers:
    ax.bar_label(containers)
#ax.bar_label(ax.containers[0])
plt.title("Churn by Payment Method")
plt.xticks(rotation=10)
plt.show()

```



**Customer is churn out when they are using Electronic Check**

```

# Numeric summary
df.describe()

   SeniorCitizen      tenure  MonthlyCharges  TotalCharges
Churn
count    7043.000000  7043.000000    7043.000000  7043.000000
7043.000000
mean     0.162147    32.371149     64.761692  2281.916928
0.265370
std      0.368612    24.559481     30.090047  2265.270398
0.441561
min     0.000000    0.000000     18.250000  18.800000
0.000000
25%     0.000000    9.000000     35.500000  402.225000
0.000000
50%     0.000000   29.000000     70.350000 1397.475000
0.000000

```

```


75%          0.000000    55.000000    89.850000  3786.600000
1.000000
max         1.000000    72.000000   118.750000  8684.800000
1.000000

```

```

# Categorical distribution example
plt.figure(figsize=(8,4))
ax=sns.countplot(x='Contract',hue='Churn',data=df)
for containers in ax.containers:
    ax.bar_label(containers)
#ax.bar_label(ax.containers[0])
plt.title("Churn by Contract Type")
plt.show()

```



in above you can say month to month contract base customer is churn more and for long term one year and 2 year customer is stay. for that we need to convey to customer contract for long time

### Feature Encoding & Scaling

```

# Encode categorical columns(in this categorical columns store in 1 variable)
cat_col=df.select_dtypes(include='object').columns
cat_col

Index(['gender', 'Partner', 'Dependents', 'PhoneService',
       'MultipleLines',
       'InternetService', 'OnlineSecurity', 'OnlineBackup',
       'DeviceProtection',

```

```

'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
'PaperlessBilling', 'PaymentMethod'],
dtype='object')

#in this we convert categorical columns in True False Format and drop
firt original column df =
df =pd.get_dummies(df,columns=cat_col,drop_first=True)

[col for col in x.columns if 'Churn' in col]

[]

# Split data in x and y
x = pd.get_dummies(df.drop('Churn', axis=1), drop_first=True)
y = df['Churn']

x.columns

Index(['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges',
       'gender_Male', 'Partner_Yes', 'Dependents_Yes',
       'PhoneService_Yes',
       'MultipleLines_No phone service', 'MultipleLines_Yes',
       'InternetService_Fiber optic', 'InternetService_No',
       'OnlineSecurity_No internet service', 'OnlineSecurity_Yes',
       'OnlineBackup_No internet service', 'OnlineBackup_Yes',
       'DeviceProtection_No internet service', 'DeviceProtection_Yes',
       'TechSupport_No internet service', 'TechSupport_Yes',
       'StreamingTV_No internet service', 'StreamingTV_Yes',
       'StreamingMovies_No internet service', 'StreamingMovies_Yes',
       'Contract_One year', 'Contract_Two year',
       'PaperlessBilling_Yes',
       'PaymentMethod_Credit card (automatic)',
       'PaymentMethod_Electronic check', 'PaymentMethod_Mailed
check'],
      dtype='object')

# Train-test split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,rando
m_state=42,stratify=y)
'''random_state : default=None
   Controls the shuffling applied to the data before applying the
split.
shuffle : default=True
   Whether or not to shuffle the data before splitting. If
shuffle=False
   then stratify must be None.'''
'random_state : default=None\n    Controls the shuffling applied to
the data before applying the split.\nshuffle : default=True\nWhether or not to shuffle the data before splitting. If shuffle=False\
n    then stratify must be None.'

```

```
# Scale
scaler=StandardScaler()
x_train=scaler.fit_transform(x_train)
x_test=scaler.transform(x_test)
```

### Model Training (Logistic Regression & Random Forest)

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

models={'Logistic Regression':LogisticRegression(),
        'Random
Forest':RandomForestClassifier(n_estimators=100,random_state=42)}
'''n_estimators : int, default=100
   The number of trees in the forest.
random_state :default=None
   Controls both the randomness of the bootstrapping of the samples
used
   when building trees (if ``bootstrap=True``) and the sampling of
the
   features to consider when looking for the best split at each
node'''

'n_estimators : int, default=100\n    The number of trees in the
forest.\nrandom_state :default=None\n    Controls both the randomness
of the bootstrapping of the samples used\n    when building trees (if
``bootstrap=True``) and the sampling of the\n    features to consider
when looking for the best split at each node'

for name,model in models.items():
    model.fit(x_train,y_train)
    y_pred=model.predict(x_test)
    print(f"== {name} ===")
    print('accuracy score',accuracy_score(y_test,y_pred))
    print("ROC-AUC:", roc_auc_score(y_test, y_pred))
    print("Classification Report:\n", classification_report(y_test,
y_pred))
    print("-" * 40)

==== Logistic Regression ===
accuracy score 0.8069552874378992
ROC-AUC: 0.7302823632746905
Classification Report:
             precision      recall   f1-score   support
              0          0.85       0.89       0.87      1035
              1          0.66       0.57       0.61       374
      accuracy           0.75       0.73       0.74      1409
     macro avg            0.75       0.73       0.74      1409
```

```

weighted avg      0.80      0.81      0.80      1409

-----
== Random Forest ==
accuracy score 0.7856635911994322
ROC-AUC: 0.6918830246195975
Classification Report:
precision    recall   f1-score   support
0            0.83    0.89    0.86    1035
1            0.62    0.49    0.55     374
accuracy          0.79      -        -    1409
macro avg       0.73    0.69    0.70    1409
weighted avg    0.77    0.79    0.78    1409

```

in above Logistic Regression outperforms Random Forest in this case, with higher accuracy, ROC-AUC, and recall. Therfore Logistic Regression It's the better choice for predicting churn in this dataset.

#### Feature Importance (for Random Forest)

```

x.columns
Index(['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges',
       'gender_Male', 'Partner_Yes', 'Dependents_Yes',
       'PhoneService_Yes',
       'MultipleLines_No phone service', 'MultipleLines_Yes',
       'InternetService_Fiber optic', 'InternetService_No',
       'OnlineSecurity_No internet service', 'OnlineSecurity_Yes',
       'OnlineBackup_No internet service', 'OnlineBackup_Yes',
       'DeviceProtection_No internet service', 'DeviceProtection_Yes',
       'TechSupport_No internet service', 'TechSupport_Yes',
       'StreamingTV_No internet service', 'StreamingTV_Yes',
       'StreamingMovies_No internet service', 'StreamingMovies_Yes',
       'Contract_One year', 'Contract_Two year',
       'PaperlessBilling_Yes',
       'PaymentMethod_Credit card (automatic)',
       'PaymentMethod_Electronic check', 'PaymentMethod_Mailed
check'],
      dtype='object')

rf = RandomForestClassifier(random_state=42)
rf.fit(x, y)

RandomForestClassifier(random_state=42)

```


```

importances = pd.Series(rf.feature_importances_, index=x.columns)
importances.sort_values(ascending=False).head(10)

TotalCharges          0.192846
tenure                 0.173395
MonthlyCharges         0.172255
InternetService_Fiber optic 0.039823
PaymentMethod_Electronic check 0.035774
Contract_Two year      0.032939
gender_Male              0.028368
OnlineSecurity_Yes       0.026899
PaperlessBilling_Yes     0.025229
TechSupport_Yes           0.023075
dtype: float64

plt.figure(figsize=(8,5))
sns.barplot(
    x=importances.nlargest(10),
    y=importances.nlargest(10).index,
    palette='Blues_r'
)
plt.title("Top 10 Features Influencing Customer Churn")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.show()

```



**Financial factors (charges) and customer tenure are the strongest churn drivers. Service experience (security, support) helps retain customers. Payment and internet type also influence behavior, showing that both pricing and service quality matter. In short: High charges + short tenure + fiber optic users + electronic check payments = higher churn risk**

```
df.groupby('gender_Male')['Churn'].value_counts(normalize=True)

gender_Male  Churn
False        0      0.730791
              1      0.269209
True         0      0.738397
              1      0.261603
Name: proportion, dtype: float64

[col for col in df.columns if 'Churn' in col]

['Churn']
```